

# **Flu Shot Learning: Predict Seasonal Flu Vaccines**

**Main objective of the analysis that specifies whether your model will be focused on prediction or interpretation**

The project is focus solely on prediction of flu vaccines.

## **Brief description of the data set and a summary of its attributes**

The data for this competition comes from the National 2009 H1N1 Flu Survey (NHFS).

The National 2009 H1N1 Flu Survey (NHFS) was sponsored by the National Center for Immunization and Respiratory Diseases (NCIRD) and conducted jointly by NCIRD and the National Center for Health Statistics (NCHS), Centers for Disease Control and Prevention (CDC). The NHFS was a list-assisted random-digit-dialing telephone survey of households, designed to monitor influenza immunization coverage in the 2009-10 season.

The target population for the NHFS was all persons 6 months or older living in the United States at the time of the interview. Data from the NHFS were used to produce timely estimates of vaccination coverage rates for both the monovalent pH1N1 and trivalent seasonal influenza vaccines.

The NHFS was conducted between October 2009 and June 2010. It was one-time survey designed specifically to monitor vaccination during the 2009-2010 flu season in response to the 2009 H1N1 pandemic. The CDC has other ongoing programs for annual phone surveys that continue to monitor seasonal flu vaccination.

## **Problem description**

Can you predict whether people got seasonal flu vaccines using information they shared about their backgrounds, opinions, and health behaviors?

In this challenge, we will take a look at vaccination, a key public health measure used to fight infectious diseases. Vaccines provide immunization for individuals, and enough immunization in a community can further reduce the spread of diseases through "herd immunity."

## **The features in this dataset**

Your goal is to predict how likely individuals are to receive their seasonal flu vaccines.

Field	Description
seasonal_vaccine	Whether respondent received seasonal flu vaccine
respondent_id	a unique and random identifier
behavioral_antiviral_meds	Has taken antiviral medications. (binary)
behavioral_avoidance	Has avoided close contact with others with flu-like symptoms. (binary)
behavioral_face_mask	Has bought a face mask. (binary)
behavioral_wash_hands	Has frequently washed hands or used hand sanitizer. (binary)
behavioral_large_gatherings	Has reduced time at large gatherings. (binary)
behavioral_outside_home	Has reduced contact with people outside of own household. (binary)
behavioral_touch_face	Has avoided touching eyes, nose, or mouth. (binary)
doctor_recc_seasonal	Seasonal flu vaccine was recommended by doctor. (binary)
chronic_med_condition	Has any of the following chronic medical conditions: asthma or an other lung condition, diabetes, a heart condition, a kidney condition, sickle cell anemia or other anemia, a neurological or neuromuscular condition, a liver condition, or a weakened immune system caused by a chronic illness or by medicines taken for a chronic illness. (binary)
child_under_6_months	Has regular close contact with a child under the age of six months. (binary)
health_worker	Is a healthcare worker. (binary)
health_insurance	Has health insurance. (binary)
opinion_seas_vacc_effective	Respondent's opinion about seasonal flu vaccine effectiveness. 1 = Not at all effective; 2 = Not very effective; 3 = Don't know; 4 = Somewhat effective; 5 = Very effective.
opinion_seas_risk	Respondent's opinion about risk of getting sick with seasonal flu without vaccine. 1 = Very Low; 2 = Somewhat low; 3 = Don't know; 4 = Somewhat high; 5 = Very high.
opinion_seas_sick_from_vacc	Respondent's worry of getting sick from taking seasonal flu vaccine. 1 = Not at all worried; 2 = Not very worried; 3 = Don't know; 4 = Somewhat worried; 5 = Very worried.
age_group	Age group of respondent.
education	Self-reported education level.
race	Race of respondent.
sex	Sex of respondent.
income_poverty	Household annual income of respondent with respect to 2008 Census poverty thresholds.
marital_status	Marital status of respondent.
employment_status	Employment status of respondent.
hhs_geo_region	Respondent's residence using a 10-region geographic classification defined by the U.S. Dept. of Health and Human Services. Values are represented as short random character strings.
census_msa	Respondent's residence within metropolitan statistical areas (MSA) as defined by the U.S. Census.
household_adults	Number of other adults in household, top-coded to 3.
household_children	Number of children in household, top-coded to 3.
employment_industry	Type of industry respondent is employed in. Values are represented as short random character strings.
employment_occupation	Type of occupation of respondent. Values are represented as short random character strings.

## Brief summary of data exploration and actions taken for data cleaning and feature engineering

Data Exploration includes data summary, statistics, relevant graphs to find any relationships within.

As for data cleaning, we will check for missing values and decide what imputation method. We also check for data duplicates and outliers. Finally perform binary encoding before model training.

## Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier

from sklearn.model_selection import cross_val_score, train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler, OneHotEncoder, PolynomialFeatures
from sklearn.metrics import confusion_matrix, classification_report, mean_absolute_error, mean_squared_error, r2_score
from sklearn.metrics import plot_confusion_matrix, plot_precision_recall_curve, plot_roc_curve, accuracy_score
from sklearn.metrics import auc, f1_score, precision_score, recall_score, roc_auc_score

%matplotlib inline
sns.set_style('dark')
sns.set(font_scale=1.2)

import warnings
warnings.filterwarnings('ignore')
import pandas.util.testing as tm
from pycaret.classification import *

np.random.seed(123)

pd.options.display.max_columns= None
#pd.options.display.max_rows = None
```

```
C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tools\_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
  import pandas.util.testing as tm
```

```
In [2]: df = pd.read_csv("training_set_features.csv")
```

```
In [3]: df
```

Out [3]:

	respondent_id	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_
0	0	0.0	0.0	0.0	
1	1	0.0	1.0	0.0	
2	2	0.0	1.0	0.0	
3	3	0.0	1.0	0.0	
4	4	0.0	1.0	0.0	
...	...	...	...	...	
26702	26702	0.0	1.0	0.0	
26703	26703	0.0	1.0	0.0	
26704	26704	0.0	1.0	1.0	
26705	26705	0.0	0.0	0.0	
26706	26706	0.0	1.0	0.0	

26707 rows × 31 columns

Dataset has 31 categorical features.

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 31 columns):
 #   Column                                  Non-Null Count  Dtype
---  -
 0   respondent_id                          26707 non-null  int64
 1   behavioral_antiviral_meds              26636 non-null  float64
 2   behavioral_avoidance                   26499 non-null  float64
 3   behavioral_face_mask                   26688 non-null  float64
 4   behavioral_wash_hands                  26665 non-null  float64
 5   behavioral_large_gatherings            26620 non-null  float64
 6   behavioral_outside_home                26625 non-null  float64
 7   behavioral_touch_face                  26579 non-null  float64
 8   doctor_recc_seasonal                   24547 non-null  float64
 9   chronic_med_condition                  25736 non-null  float64
10   child_under_6_months                  25887 non-null  float64
11   health_worker                          25903 non-null  float64
12   health_insurance                       14433 non-null  float64
13   opinion_seas_vacc_effective             26245 non-null  float64
14   opinion_seas_risk                       26193 non-null  float64
15   opinion_seas_sick_from_vacc             26170 non-null  float64
16   age_group                              26707 non-null  object
17   education                              25300 non-null  object
18   race                                   26707 non-null  object
19   sex                                    26707 non-null  object
20   income_poverty                         22284 non-null  object
21   marital_status                         25299 non-null  object
22   rent_or_own                           24665 non-null  object
23   employment_status                     25244 non-null  object
24   hhs_geo_region                         26707 non-null  object
25   census_msa                             26707 non-null  object
26   household_adults                       26458 non-null  float64
27   household_children                     26458 non-null  float64
28   employment_industry                    13377 non-null  object
29   employment_occupation                  13237 non-null  object
30   seasonal_vaccine                       26707 non-null  int64
dtypes: float64(17), int64(2), object(12)
memory usage: 6.3+ MB
```

Summary of statistics below:

```
In [5]: df.describe(include='all').T
```

```
Out[5]:
```

	count	unique	top	freq	mean	std	min	25%	50%	75%
<b>respondent_id</b>	26707	NaN	NaN	NaN	13353	7709.79	0	6676.5	13353	2002
<b>behavioral_antiviral_meds</b>	26636	NaN	NaN	NaN	0.0488437	0.215545	0	0	0	
<b>behavioral_avoidance</b>	26499	NaN	NaN	NaN	0.725612	0.446214	0	0	1	
<b>behavioral_face_mask</b>	26688	NaN	NaN	NaN	0.0689823	0.253429	0	0	0	
<b>behavioral_wash_hands</b>	26665	NaN	NaN	NaN	0.825614	0.379448	0	1	1	
<b>behavioral_large_gatherings</b>	26620	NaN	NaN	NaN	0.35864	0.47961	0	0	0	
<b>behavioral_outside_home</b>	26625	NaN	NaN	NaN	0.337315	0.472802	0	0	0	
<b>behavioral_touch_face</b>	26579	NaN	NaN	NaN	0.677264	0.467531	0	0	1	
<b>doctor_recc_seasonal</b>	24547	NaN	NaN	NaN	0.329735	0.470126	0	0	0	
<b>chronic_med_condition</b>	25736	NaN	NaN	NaN	0.283261	0.450591	0	0	0	
<b>child_under_6_months</b>	25887	NaN	NaN	NaN	0.0825897	0.275266	0	0	0	
<b>health_worker</b>	25903	NaN	NaN	NaN	0.111918	0.315271	0	0	0	
<b>health_insurance</b>	14433	NaN	NaN	NaN	0.87972	0.3253	0	1	1	
<b>opinion_seas_vacc_effective</b>	26245	NaN	NaN	NaN	4.02599	1.08656	1	4	4	
<b>opinion_seas_risk</b>	26193	NaN	NaN	NaN	2.71916	1.38506	1	2	2	
<b>opinion_seas_sick_from_vacc</b>	26170	NaN	NaN	NaN	2.11811	1.33295	1	1	2	
<b>age_group</b>	26707	5	65+ Years	6843	NaN	NaN	NaN	NaN	NaN	N
<b>education</b>	25300	4	College Graduate	10097	NaN	NaN	NaN	NaN	NaN	N
<b>race</b>	26707	4	White	21222	NaN	NaN	NaN	NaN	NaN	N
<b>sex</b>	26707	2	Female	15858	NaN	NaN	NaN	NaN	NaN	N
<b>income_poverty</b>	22284	3	<=\$75,000, Above Poverty	12777	NaN	NaN	NaN	NaN	NaN	N
<b>marital_status</b>	25299	2	Married	13555	NaN	NaN	NaN	NaN	NaN	N
<b>rent_or_own</b>	24665	2	Own	18736	NaN	NaN	NaN	NaN	NaN	N
<b>employment_status</b>	25244	3	Employed	13560	NaN	NaN	NaN	NaN	NaN	N
<b>hhs_geo_region</b>	26707	10	lzgpxyt	4297	NaN	NaN	NaN	NaN	NaN	N
<b>census_msa</b>	26707	3	MSA, Not Principle City	11645	NaN	NaN	NaN	NaN	NaN	N
<b>household_adults</b>	26458	NaN	NaN	NaN	0.886499	0.753422	0	0	1	
<b>household_children</b>	26458	NaN	NaN	NaN	0.534583	0.928173	0	0	0	
<b>employment_industry</b>	13377	21	fcxhlnwr	2468	NaN	NaN	NaN	NaN	NaN	N
<b>employment_occupation</b>	13237	23	xtkaffoo	1778	NaN	NaN	NaN	NaN	NaN	N
<b>seasonal_vaccine</b>	26707	NaN	NaN	NaN	0.465608	0.498825	0	0	0	

Shape of dataset:

```
In [6]: df.shape
```

```
Out[6]: (26707, 31)
```

```
In [7]: df.columns
```

```
Out[7]: Index(['respondent_id', 'behavioral_antiviral_meds', 'behavioral_avoidance',  
              'behavioral_face_mask', 'behavioral_wash_hands',  
              'behavioral_large_gatherings', 'behavioral_outside_home',  
              'behavioral_touch_face', 'doctor_recc_seasonal',  
              'chronic_med_condition', 'child_under_6_months', 'health_worker',  
              'health_insurance', 'opinion_seas_vacc_effective', 'opinion_seas_risk',  
              'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race', 'sex',  
              'income_poverty', 'marital_status', 'rent_or_own', 'employment_status',  
              'hhs_geo_region', 'census_msa', 'household_adults',  
              'household_children', 'employment_industry', 'employment_occupation',  
              'seasonal_vaccine'],  
             dtype='object')
```

## Data Exploration

The best way is to create graphs!

## Data Visualization

The dataset is mainly discrete/categorical types.

```
In [8]: df.hist(bins=50, figsize=(20,20))

plt.suptitle('Feature Distribution', x=0.5, y=1.02, ha='center', fontsize='large')

plt.tight_layout()

plt.show();
```



Below are each visuals of the data:



```
In [10]: fig = plt.figure(figsize=(20,40))

plt.subplot(7,2,1)
plt.title("Has taken antiviral medications")
sns.countplot(df.behavioral_antiviral_meds, hue=df.seasonal_vaccine)

plt.subplot(7,2,2)
plt.title("Has avoided close contact with others with flu-like symptoms")
sns.countplot(df.behavioral_avoidance, hue=df.seasonal_vaccine)

plt.subplot(7,2,3)
plt.title("Has bought a face mask")
sns.countplot(df.behavioral_face_mask, hue=df.seasonal_vaccine)

plt.subplot(7,2,4)
plt.title("Has frequently washed hands or used hand sanitizer")
sns.countplot(df.behavioral_wash_hands, hue=df.seasonal_vaccine)

plt.subplot(7,2,5)
plt.title("Has reduced time at large gatherings")
sns.countplot(df.behavioral_large_gatherings, hue=df.seasonal_vaccine)

plt.subplot(7,2,6)
plt.title("Has reduced contact with people outside of own household")
sns.countplot(df.behavioral_outside_home, hue=df.seasonal_vaccine)

plt.subplot(7,2,7)
plt.title("Has avoided touching eyes, nose, or mouth")
sns.countplot(df.behavioral_touch_face, hue=df.seasonal_vaccine)

plt.subplot(7,2,8)
plt.title("Seasonal flu vaccine was recommended by doctor")
sns.countplot(df.doctor_recc_seasonal, hue=df.seasonal_vaccine)



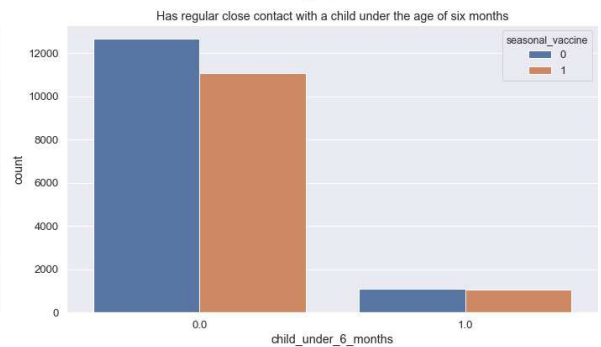

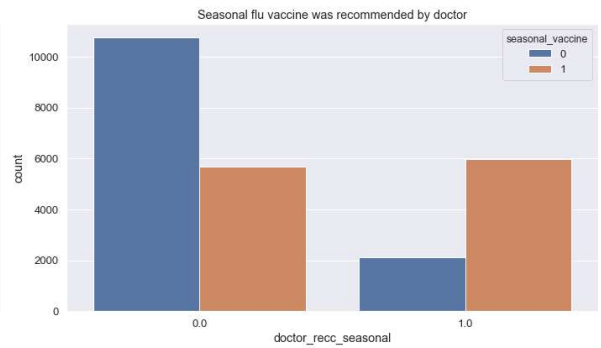
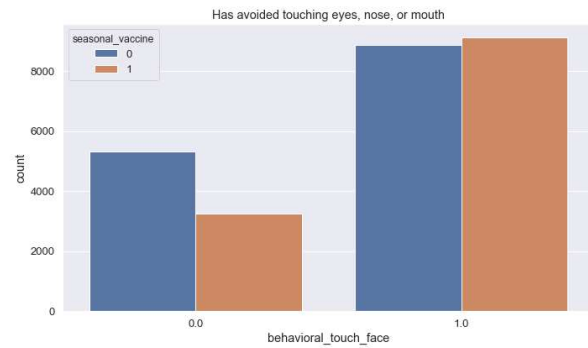


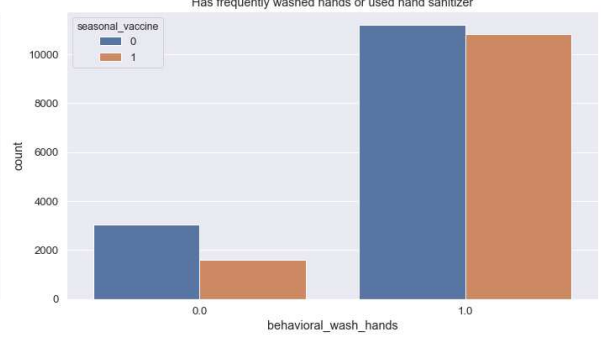
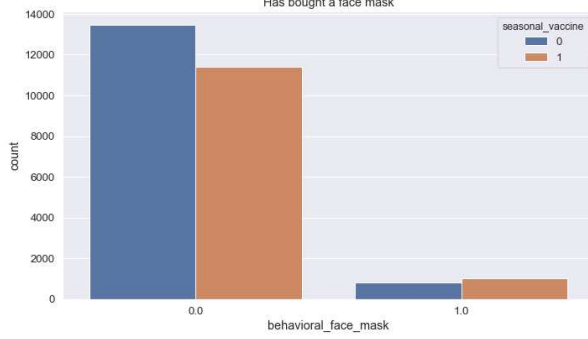
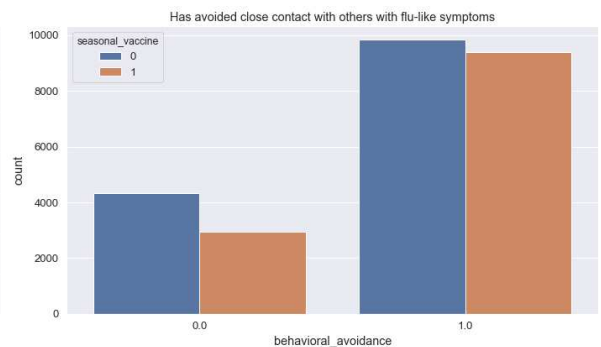

plt.subplot(7,2,9)
plt.title("Has any of the following chronic medical conditions")
sns.countplot(df.chronic_med_condition, hue=df.seasonal_vaccine)

plt.subplot(7,2,10)
plt.title("Has regular close contact with a child under the age of six months")
sns.countplot(df.child_under_6_months, hue=df.seasonal_vaccine)

plt.subplot(7,2,11)
plt.title("Is a healthcare worker")
sns.countplot(df.health_worker, hue=df.seasonal_vaccine)

plt.subplot(7,2,12)
plt.title("Has health insurance")
sns.countplot(df.health_insurance, hue=df.seasonal_vaccine)

plt.tight_layout()
plt.show()
```



**Part 1 of Data Analysis:**

Those who had vaccine avoided close contacts which is surprising since vaccines are supposed to protect them. But they didn't avoid large gatherings.

As for flu vaccine which doctor recommended, there is such acceptance among people.

Health Care workers are most vulnerable but majority of them do vaccinate.

```
In [11]: fig = plt.figure(figsize=(20,40))

plt.subplot(7,2,1)
plt.title("Respondent's opinion about seasonal flu vaccine effectiveness")
sns.countplot(df.opinion_seas_vacc_effective, hue=df.seasonal_vaccine)

plt.subplot(7,2,2)
plt.title("Respondent's opinion about risk of getting sick with seasonal flu without vaccine")
sns.countplot(df.opinion_seas_risk, hue=df.seasonal_vaccine)

plt.subplot(7,2,3)
plt.title("Respondent's worry of getting sick from taking seasonal flu vaccine")
sns.countplot(df.opinion_seas_sick_from_vacc, hue=df.seasonal_vaccine)

plt.subplot(7,2,4)
plt.title("Age group of respondent")
sns.countplot(df.age_group, hue=df.seasonal_vaccine)

plt.subplot(7,2,5)
plt.title("Self-reported education level")
sns.countplot(df.education, hue=df.seasonal_vaccine)

plt.subplot(7,2,6)
plt.title("Race of respondent")
sns.countplot(df.race, hue=df.seasonal_vaccine)

plt.subplot(7,2,7)
plt.title("Sex of respondent")
sns.countplot(df.sex, hue=df.seasonal_vaccine)

plt.subplot(7,2,8)
plt.title("Household annual income of respondent")
sns.countplot(df.income_poverty, hue=df.seasonal_vaccine)




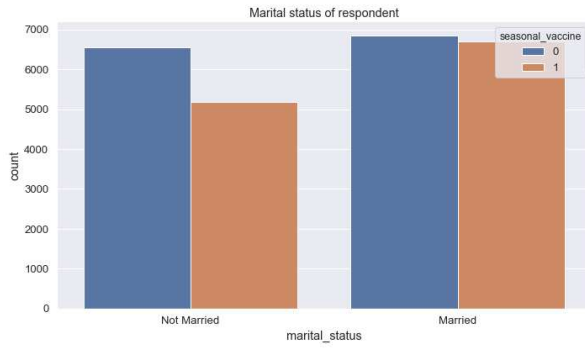

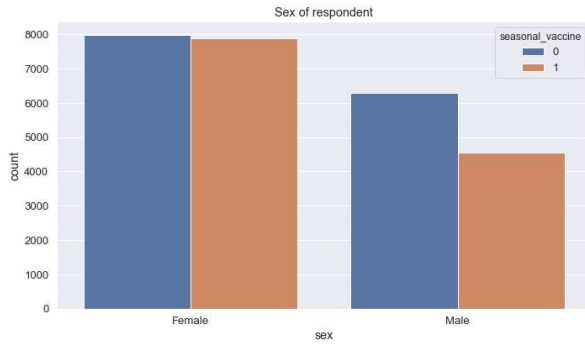
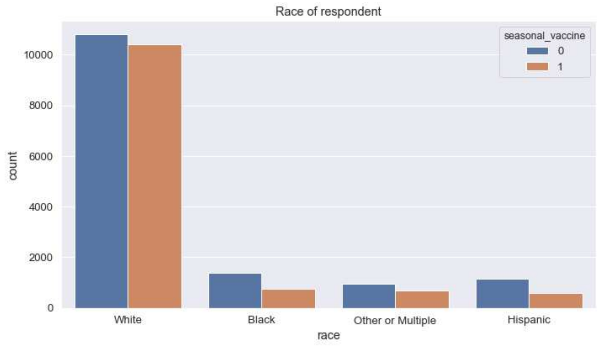


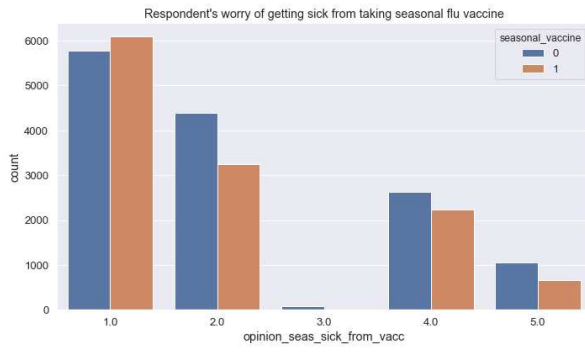


plt.subplot(7,2,9)
plt.title("Marital status of respondent")
sns.countplot(df.marital_status, hue=df.seasonal_vaccine)

plt.subplot(7,2,10)
plt.title("Housing situation of respondent")
sns.countplot(df.rent_or_own, hue=df.seasonal_vaccine)

plt.subplot(7,2,11)
plt.title("Employment status of respondent")
sns.countplot(df.employment_status, hue=df.seasonal_vaccine)

plt.subplot(7,2,12)
plt.title("Respondent's residence using a 10-region")
sns.countplot(df.hhs_geo_region, hue=df.seasonal_vaccine)

plt.tight_layout()
plt.show()
```



**Part 2 of Data Analysis:**

As for respondents opinion, risks and worry, there are no surprises for those who trust flu vaccines.

Respondents more than age 65 and College Educated are vaccinated.

Mainly whites, female, more than 75k income, married, own a house and employed can afford flu vaccines.

```

In [12]: fig = plt.figure(figsize=(20,40))

plt.subplot(7,2,1)
plt.title("Respondent's residence within metropolitan statistical areas")
sns.countplot(df.census_msa, hue=df.seasonal_vaccine)

plt.subplot(7,2,2)
plt.title("Number of other adults in household")
sns.countplot(df.household_adults, hue=df.seasonal_vaccine)

plt.subplot(7,2,3)
plt.title("Number of children in household")
sns.countplot(df.household_children, hue=df.seasonal_vaccine)

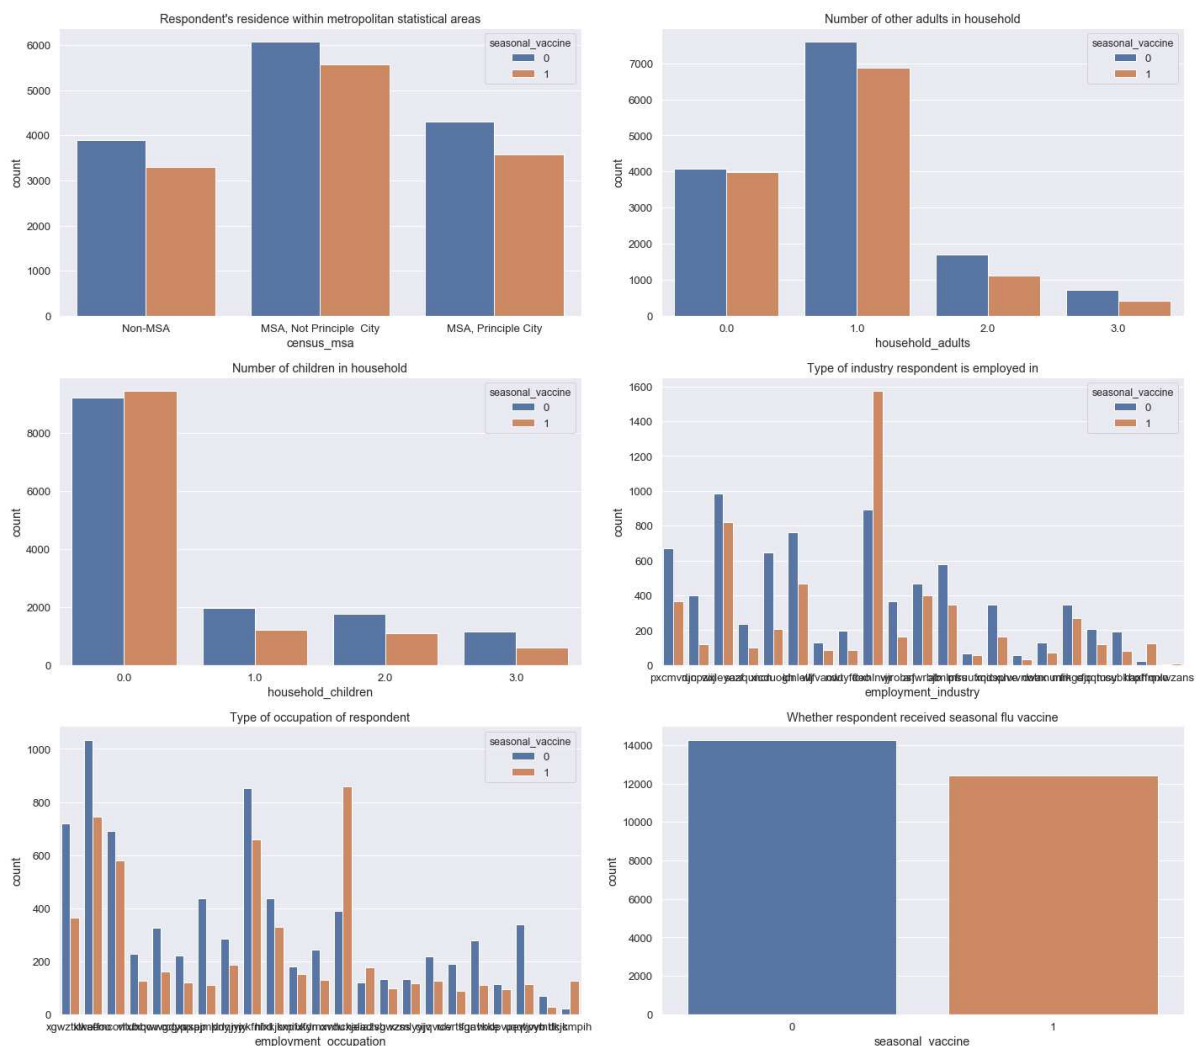
plt.subplot(7,2,4)
plt.title("Type of industry respondent is employed in")
sns.countplot(df.employment_industry, hue=df.seasonal_vaccine)

plt.subplot(7,2,5)
plt.title("Type of occupation of respondent")
sns.countplot(df.employment_occupation, hue=df.seasonal_vaccine)

plt.subplot(7,2,6)
plt.title("Whether respondent received seasonal flu vaccine")
sns.countplot(df.seasonal_vaccine)

plt.tight_layout()
plt.show()

```



### Part 3 of Data Analysis:

City dwellers, one household adults and no children mainly are vaccinated.

Unknown employment industry and occupation type is masked/not revealed to us.

As for seasonal vaccine, both are more or less equal quantity.

```
In [13]: df['seasonal_vaccine'].value_counts()
```

```
Out[13]: 0    14272
          1    12435
          Name: seasonal_vaccine, dtype: int64
```

Now we check any correlation between features:


```
In [14]: df.corr()
```

```
Out[14]:
```

	respondent_id	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_m
respondent_id	1.000000	-0.008475	0.009638	-0.00
behavioral_antiviral_meds	-0.008475	1.000000	0.049247	0.14
behavioral_avoidance	0.009638	0.049247	1.000000	0.06
behavioral_face_mask	-0.006644	0.146261	0.064946	1.00
behavioral_wash_hands	0.011105	0.064119	0.338130	0.08
behavioral_large_gatherings	0.004539	0.106287	0.227675	0.18
behavioral_outside_home	0.009011	0.127679	0.220348	0.16
behavioral_touch_face	0.007575	0.070868	0.335335	0.10
doctor_recc_seasonal	0.001500	0.030909	0.074088	0.06
chronic_med_condition	0.005797	0.008465	0.039435	0.06
child_under_6_months	-0.004839	0.028788	-0.000414	0.03
health_worker	-0.003149	0.009465	0.001180	0.06
health_insurance	-0.012603	-0.063988	0.032662	-0.04
opinion_seas_vacc_effective	0.005935	0.015003	0.119554	0.04
opinion_seas_risk	-0.005291	0.085315	0.129504	0.11
opinion_seas_sick_from_vacc	0.009563	0.084305	0.082942	0.09
household_adults	0.000187	0.044900	0.019122	0.01
household_children	-0.003726	0.084822	0.040328	0.00
seasonal_vaccine	-0.004652	0.006277	0.076395	0.05



```
In [15]: plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), cmap="coolwarm", annot=True, fmt='.2f', linewidths=2)
plt.show()
```



## Quote:

"Factors that may bias the results of observational studies can be broadly categorized as: selection bias resulting from the way study subjects are recruited or from differing rates of study participation depending on the subjects' cultural background, age, or socioeconomic status, information bias, measurement error, confounders, and further factors."

We will drop a number of features which we think that will make the model biased to a certain group/gender/income/social.

## Drop unwanted features

```
In [16]: df.columns
```

```
Out[16]: Index(['respondent_id', 'behavioral_antiviral_meds', 'behavioral_avoidance',
               'behavioral_face_mask', 'behavioral_wash_hands',
               'behavioral_large_gatherings', 'behavioral_outside_home',
               'behavioral_touch_face', 'doctor_recc_seasonal',
               'chronic_med_condition', 'child_under_6_months', 'health_worker',
               'health_insurance', 'opinion_seas_vacc_effective', 'opinion_seas_risk',
               'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race', 'sex',
               'income_poverty', 'marital_status', 'rent_or_own', 'employment_status',
               'hhs_geo_region', 'census_msa', 'household_adults',
               'household_children', 'employment_industry', 'employment_occupation',
               'seasonal_vaccine'],
              dtype='object')
```

```
In [17]: df.drop(['respondent_id', 'health_insurance', 'age_group', 'education', 'race', 'sex',
                  'income_poverty', 'marital_status', 'rent_or_own',
                  'employment_status', 'hhs_geo_region', 'census_msa', 'household_adults', 'household_children',
                  'employment_industry', 'employment_occupation'], axis=1, inplace=True)
```


```
In [18]: df
```

Out [18]:

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0.0	0.0	0.0	0.0	0.0
1	0.0	1.0	0.0	0.0	1.0
2	0.0	1.0	0.0	0.0	0.0
3	0.0	1.0	0.0	0.0	1.0
4	0.0	1.0	0.0	0.0	1.0
...	...	...	...	...	...
26702	0.0	1.0	0.0	0.0	0.0
26703	0.0	1.0	0.0	0.0	1.0
26704	0.0	1.0	1.0	1.0	1.0
26705	0.0	0.0	0.0	0.0	0.0
26706	0.0	1.0	0.0	0.0	0.0

26707 rows x 15 columns

```
In [19]: sns.pairplot(df.sample(500), hue='seasonal_vaccine')
plt.show()
```



## Treat Missing Values

```
In [20]: df.isnull().sum()
```

```
Out [20]: behavioral_antiviral_meds      71
behavioral_avoidance                  208
behavioral_face_mask                  19
behavioral_wash_hands                  42
behavioral_large_gatherings            87
behavioral_outside_home                 82
behavioral_touch_face                  128
doctor_recc_seasonal                 2160
chronic_med_condition                  971
child_under_6_months                  820
health_worker                         804
opinion_seas_vacc_effective            462
opinion_seas_risk                     514
opinion_seas_sick_from_vacc           537
seasonal_vaccine                       0
dtype: int64
```

```
In [21]: df.dropna(inplace=True)
```

```
In [22]: df.isnull().sum()
```

```
Out [22]: behavioral_antiviral_meds      0
behavioral_avoidance                    0
behavioral_face_mask                    0
behavioral_wash_hands                    0
behavioral_large_gatherings              0
behavioral_outside_home                  0
behavioral_touch_face                    0
doctor_recc_seasonal                    0
chronic_med_condition                    0
child_under_6_months                    0
health_worker                           0
opinion_seas_vacc_effective              0
opinion_seas_risk                       0
opinion_seas_sick_from_vacc              0
seasonal_vaccine                        0
dtype: int64
```

```
In [23]: df.reset_index(drop=True, inplace=True)
```

```
In [24]: df
```

```
Out[24]:
```

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0.0	0.0	0.0	0.0	
1	0.0	1.0	0.0	1.0	
2	0.0	1.0	0.0	1.0	
3	0.0	1.0	0.0	1.0	
4	0.0	1.0	0.0	1.0	
...	...	...	...	...	...
23183	0.0	0.0	0.0	1.0	
23184	0.0	1.0	0.0	0.0	
23185	0.0	1.0	0.0	1.0	
23186	0.0	1.0	1.0	1.0	
23187	0.0	1.0	0.0	0.0	

23188 rows x 15 columns

```
In [25]: df['seasonal_vaccine'].value_counts()
```

```
Out[25]: 0    12111
         1    11077
         Name: seasonal_vaccine, dtype: int64
```

```
In [26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23188 entries, 0 to 23187
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   behavioral_antiviral_meds             23188 non-null  float64
1   behavioral_avoidance                  23188 non-null  float64
2   behavioral_face_mask                  23188 non-null  float64
3   behavioral_wash_hands                 23188 non-null  float64
4   behavioral_large_gatherings           23188 non-null  float64
5   behavioral_outside_home               23188 non-null  float64
6   behavioral_touch_face                 23188 non-null  float64
7   doctor_recc_seasonal                  23188 non-null  float64
8   chronic_med_condition                23188 non-null  float64
9   child_under_6_months                 23188 non-null  float64
10  health_worker                        23188 non-null  float64
11  opinion_seas_vacc_effective            23188 non-null  float64
12  opinion_seas_risk                     23188 non-null  float64
13  opinion_seas_sick_from_vacc           23188 non-null  float64
14  seasonal_vaccine                     23188 non-null  int64
dtypes: float64(14), int64(1)
memory usage: 2.7 MB
```

```
In [27]: df = df.astype('int8') #Change to integer type
```

```
In [28]: df.dtypes
```

```
Out[28]: behavioral_antiviral_meds      int8
behavioral_avoidance                  int8
behavioral_face_mask                  int8
behavioral_wash_hands                  int8
behavioral_large_gatherings            int8
behavioral_outside_home                int8
behavioral_touch_face                  int8
doctor_recc_seasonal                  int8
chronic_med_condition                  int8
child_under_6_months                  int8
health_worker                         int8
opinion_seas_vacc_effective            int8
opinion_seas_risk                     int8
opinion_seas_sick_from_vacc            int8
seasonal_vaccine                      int8
dtype: object
```

```
In [29]: df
```

```
Out[29]:
```

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0	0	0	0	
1	0	1	0	1	
2	0	1	0	1	
3	0	1	0	1	
4	0	1	0	1	
...	...	...	...	...	...
23183	0	0	0	1	
23184	0	1	0	0	
23185	0	1	0	1	
23186	0	1	1	1	
23187	0	1	0	0	

23188 rows × 5 columns

```
In [30]: df.describe()
```

```
Out[30]:
```

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
count	23188.000000	23188.000000	23188.000000	23188.000000	
mean	0.049336	0.731197	0.068139	0.829481	
std	0.216573	0.443347	0.251989	0.376096	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	1.000000	
50%	0.000000	1.000000	0.000000	1.000000	
75%	0.000000	1.000000	0.000000	1.000000	
max	1.000000	1.000000	1.000000	1.000000	

```
In [31]: df['opinion_seas_vacc_effective'] = df['opinion_seas_vacc_effective'].astype('object')
```

```
In [32]: df['opinion_seas_risk'] = df['opinion_seas_risk'].astype('object')
```

```
In [33]: df['opinion_seas_sick_from_vacc'] = df['opinion_seas_sick_from_vacc'].astype('object')
```

```
In [34]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 23188 entries, 0 to 23187
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   behavioral_antiviral_meds             23188 non-null   int8
1   behavioral_avoidance                  23188 non-null   int8
2   behavioral_face_mask                  23188 non-null   int8
3   behavioral_wash_hands                 23188 non-null   int8
4   behavioral_large_gatherings           23188 non-null   int8
5   behavioral_outside_home               23188 non-null   int8
6   behavioral_touch_face                 23188 non-null   int8
7   doctor_recc_seasonal                  23188 non-null   int8
8   chronic_med_condition                 23188 non-null   int8
9   child_under_6_months                 23188 non-null   int8
10  health_worker                         23188 non-null   int8
11  opinion_seas_vacc_effective            23188 non-null   object
12  opinion_seas_risk                      23188 non-null   object
13  opinion_seas_sick_from_vacc            23188 non-null   object
14  seasonal_vaccine                      23188 non-null   int8
dtypes: int8(12), object(3)
memory usage: 815.3+ KB
```

## Create dummy variables

```
In [35]: df2 = pd.get_dummies(data=df, drop_first=True)
```

```
In [36]: df2
```

Out [36]:

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0	0	0	0	
1	0	1	0	1	
2	0	1	0	1	
3	0	1	0	1	
4	0	1	0	1	
...	...	...	...	...	...
23183	0	0	0	1	
23184	0	1	0	0	
23185	0	1	0	1	
23186	0	1	1	1	
23187	0	1	0	0	

23188 rows × 24 columns

```
In [37]: df2.columns
```

```
Out[37]: Index(['behavioral_antiviral_meds', 'behavioral_avoidance',  
              'behavioral_face_mask', 'behavioral_wash_hands',  
              'behavioral_large_gatherings', 'behavioral_outside_home',  
              'behavioral_touch_face', 'doctor_recc_seasonal',  
              'chronic_med_condition', 'child_under_6_months', 'health_worker',  
              'seasonal_vaccine', 'opinion_seas_vacc_effective_2',  
              'opinion_seas_vacc_effective_3', 'opinion_seas_vacc_effective_4',  
              'opinion_seas_vacc_effective_5', 'opinion_seas_risk_2',  
              'opinion_seas_risk_3', 'opinion_seas_risk_4', 'opinion_seas_risk_5',  
              'opinion_seas_sick_from_vacc_2', 'opinion_seas_sick_from_vacc_3',  
              'opinion_seas_sick_from_vacc_4', 'opinion_seas_sick_from_vacc_5'],  
              dtype='object')
```

```
In [38]: df2 = df2[['behavioral_antiviral_meds', 'behavioral_avoidance',  
                  'behavioral_face_mask', 'behavioral_wash_hands',  
                  'behavioral_large_gatherings', 'behavioral_outside_home',  
                  'behavioral_touch_face', 'doctor_recc_seasonal',  
                  'chronic_med_condition', 'child_under_6_months', 'health_worker',  
                  'opinion_seas_vacc_effective_2',  
                  'opinion_seas_vacc_effective_3', 'opinion_seas_vacc_effective_4',  
                  'opinion_seas_vacc_effective_5', 'opinion_seas_risk_2',  
                  'opinion_seas_risk_3', 'opinion_seas_risk_4', 'opinion_seas_risk_5',  
                  'opinion_seas_sick_from_vacc_2', 'opinion_seas_sick_from_vacc_3',  
                  'opinion_seas_sick_from_vacc_4', 'opinion_seas_sick_from_vacc_5', 'seasonal_vaccine']]
```

```
In [39]: df2
```

```
Out[39]:
```

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0	0	0	0	
1	0	1	0	1	
2	0	1	0	1	
3	0	1	0	1	
4	0	1	0	1	
...	...	...	...	...	...
23183	0	0	0	1	
23184	0	1	0	0	
23185	0	1	0	1	
23186	0	1	1	1	
23187	0	1	0	0	

23188 rows × 24 columns

## Create and save processed dataset

```
In [40]: df2.to_csv("train.csv", index=False)
```

```
In [ ]:
```

```
In [41]: df = pd.read_csv("train.csv")
```

```
In [42]: df
```

```
Out[42]:
```

	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0	0	0	0	0
1	0	1	0	1	1
2	0	1	0	1	1
3	0	1	0	1	1
4	0	1	0	1	1
...	...	...	...	...	...
23183	0	0	0	1	1
23184	0	1	0	0	0
23185	0	1	0	1	1
23186	0	1	1	1	1
23187	0	1	0	0	0

23188 rows × 24 columns

```
In [43]: df.shape
```

```
Out[43]: (23188, 24)
```

**Summary of training at least three different classifier models, preferably of different nature in explainability and predictability. For example, you can start with a simple logistic regression as a baseline, adding other models or ensemble models. Preferably, all your models use the same training and test splits, or the same cross-validation method.**

## Train Test Split

```
In [44]: X = df.iloc[:,0:23]
y = df.iloc[:,23]
```

```
In [45]: X.values
```

```
Out[45]: array([[0, 0, 0, ..., 0, 0, 0],
                [0, 1, 0, ..., 0, 1, 0],
                [0, 1, 0, ..., 0, 0, 0],
                ...,
                [0, 1, 0, ..., 0, 0, 0],
                [0, 1, 1, ..., 0, 0, 0],
                [0, 1, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [46]: y.values
```

```
Out[46]: array([0, 1, 1, ..., 0, 1, 0], dtype=int64)
```

```
In [47]: X_train, X_test, y_train, y_test = train_test_split(X.values, y.values, test_size=
0.2, random_state=123, stratify=y)
```



```
In [48]: X_train
```

```
Out[48]: array([[0, 1, 0, ..., 0, 0, 0],
               [0, 1, 0, ..., 0, 0, 0],
               [0, 1, 0, ..., 0, 0, 0],
               ...,
               [0, 1, 1, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0],
               [0, 1, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [49]: X_test
```

```
Out[49]: array([[0, 1, 0, ..., 0, 0, 0],
               [0, 1, 0, ..., 0, 0, 0],
               [0, 0, 0, ..., 0, 0, 0],
               ...,
               [0, 0, 0, ..., 0, 0, 0],
               [1, 0, 0, ..., 0, 0, 0],
               [0, 1, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [50]: y_train
```

```
Out[50]: array([1, 0, 1, ..., 1, 1, 1], dtype=int64)
```

## Logistic Regression

```
In [51]: lr = LogisticRegression(random_state=123)
```

```
In [52]: lr.fit(X_train, y_train)
```

```
Out[52]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=123, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [53]: lr.coef_
```

```
Out[53]: array([[ -0.23062336,  -0.07794305,  -0.00527337,   0.09815467,   0.00333909,
                  -0.07933634,   0.27948713,   1.36430408,   0.35155097,  -0.16748045,
                   0.79348647,  -0.29698112,   0.6244859 ,   0.75495696,   1.77721418,
                   0.80614637,   1.6792209 ,   1.70246871,   2.0063813 ,  -0.44950573,
                  -1.62206222,  -0.67756796,  -1.24302248]])
```

```
In [54]: lr.intercept_
```

```
Out[54]: array([-2.63615376])
```

```
In [55]: ypred_lr = lr.predict(X_test)
```

```
In [56]: y_test[:10]
```

```
Out[56]: array([1, 0, 1, 0, 1, 1, 1, 0, 1, 1], dtype=int64)
```

```
In [57]: ypred_lr[:10]
```


```
Out[57]: array([1, 0, 0, 0, 1, 1, 1, 0, 0, 1], dtype=int64)
```

## Logistic Regression Model Evaluation

```
In [58]: cm = confusion_matrix(y_test,ypred_lr)
cm
```

```
Out[58]: array([[1935,  487],
               [ 577, 1639]], dtype=int64)
```


```
In [59]: fig , ax = plt.subplots(figsize=(10,5))
sns.heatmap(cm, annot=True,fmt='.4g',linewidths=2, cmap='viridis')
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



```
In [60]: print(classification_report(y_test,ypred_lr))
```

	precision	recall	f1-score	support
0	0.77	0.80	0.78	2422
1	0.77	0.74	0.75	2216
accuracy			0.77	4638
macro avg	0.77	0.77	0.77	4638
weighted avg	0.77	0.77	0.77	4638

```
In [61]: plot_roc_curve(lr,X_test,y_test)
plt.show()
```



```
In [62]: accuracy_score(y_test,ypred_lr)
```

```
Out[62]: 0.7705907718844329
```

## Random Forest Classifier

```
In [63]: rf = RandomForestClassifier(random_state=123)
```

```
In [64]: rf.fit(X_train, y_train)
```

```
Out[64]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=123,
                                verbose=0, warm_start=False)
```

```
In [65]: ypred_rf = rf.predict(X_test)
```

```
In [66]: y_test[:10]
```

```
Out[66]: array([1, 0, 1, 0, 1, 1, 1, 0, 1, 1], dtype=int64)
```

```
In [67]: ypred_rf[:10]
```


```
Out[67]: array([1, 0, 0, 0, 1, 1, 1, 1, 0, 1], dtype=int64)
```

## Random Forest Model Evaluation

```
In [68]: cm = confusion_matrix(y_test,ypred_rf)
cm
```

```
Out[68]: array([[1827,  595],
                [ 638, 1578]], dtype=int64)
```


```
In [69]: fig, ax = plt.subplots(figsize=(10,5))
sns.heatmap(cm, annot=True,fmt='.4g',linewidths=2, cmap='viridis')
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



```
In [70]: print(classification_report(y_test,ypred_rf))
```

	precision	recall	f1-score	support
0	0.74	0.75	0.75	2422
1	0.73	0.71	0.72	2216
accuracy			0.73	4638
macro avg	0.73	0.73	0.73	4638
weighted avg	0.73	0.73	0.73	4638

```
In [71]: plot_roc_curve(rf,X_test,y_test)
plt.show()
```



```
In [72]: accuracy_score(y_test,ypred_rf)
```

```
Out[72]: 0.7341526520051747
```

## Gradient Boosting Classifier

```
In [73]: gbc = GradientBoostingClassifier(random_state=123)
```

```
In [74]: gbc.fit(X_train,y_train)
```

```
Out[74]: GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
    learning_rate=0.1, loss='deviance', max_depth=3,
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    n_iter_no_change=None, presort='deprecated',
    random_state=123, subsample=1.0, tol=0.0001,
    validation_fraction=0.1, verbose=0,
    warm_start=False)
```

```
In [75]: ypredgbc = gbc.predict(X_test)
```

```
In [76]: y_test[:10]
```

```
Out[76]: array([1, 0, 1, 0, 1, 1, 1, 0, 1, 1], dtype=int64)
```

```
In [77]: ypredgbc[:10]
```


```
Out[77]: array([1, 0, 0, 0, 1, 1, 1, 1, 0, 1], dtype=int64)
```

## Gradient Boosting Model Evaluation

```
In [78]: cm = confusion_matrix(y_test,ypredgbc)
cm
```

```
Out[78]: array([[1922,  500],
               [ 575, 1641]], dtype=int64)
```


```
In [79]: fig , ax = plt.subplots(figsize=(10,5))
sns.heatmap(cm, annot=True,fmt='.4g',linewidths=2, cmap='viridis')
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



```
In [80]: print(classification_report(y_test,ypredgbc))
```

	precision	recall	f1-score	support
0	0.77	0.79	0.78	2422
1	0.77	0.74	0.75	2216
accuracy			0.77	4638
macro avg	0.77	0.77	0.77	4638
weighted avg	0.77	0.77	0.77	4638

```
In [81]: plot_roc_curve(gbc,X_test,y_test)
plt.show()
```



```
In [82]: accuracy_score(y_test,ypredgbc)
```

```
Out[82]: 0.7682190599396291
```

**A paragraph explaining which of your classifier models you recommend as a final model that best fits your needs in terms of accuracy and explainability.**

Logistic Regression gives us the best accuracy and F1 score. Therefore it is recommended.

**Summary Key Findings and Insights, which walks your reader through the main drivers of your model and insights from your data derived from your classifier model.**

The features we selected gave us a decent accuracy and good result. The result differences are small and we select Logistic Regression because it's a simple model.

**Suggestions for next steps in analyzing this data, which may include suggesting revisiting this model adding specific data features to achieve a better explanation or a better prediction.**

For features that are biased, we need to gather more data and made equal values for race, sex, income etc. We have to ensure the model we developed stays bias free.

We can also explore other models like decision tree, support vector machine, KNN classifiers model to see if they can able to analyse the data patterns to give better predictions. We also can adjust hyperparameters for each model to get better results.